

Evaluation of Indic Translation Models

Task 1 Project Report · Parts A, B & C
Gourilakshmi S

Overview

This project evaluates five translation models on English-to-Tamil conversion across three analytical lenses: end-to-end translation quality via sacreBLEU, chrF, chrF++, TER, and ROUGE-L (Part A), token-level comparative metrics across all five models (Part B), and a deep-dive into how each tokenizer handles raw Tamil Unicode script without any translation (Part C). Together, the three parts build a single cohesive argument: tokenizer vocabulary design is the decisive factor in Tamil NLP quality, and the numbers at every level prove it.

Dataset (Part A)	500 English-Tamil pairs · Helsinki-NLP/opus-100 en-ta test split
Sentences (Parts B & C)	15 English (3 complexity tiers) · 20 raw Tamil sentences
Primary Metric	sacreBLEU, chrF, chrF++, TER, ROUGE-L (Part A) · Expansion, fragmentation, UNK rate (B/C)
Framework	HuggingFace Transformers · IndicTransToolkit · SentencePiece

Models Evaluated

Model	Model ID	Tamil Support	Expansion
IndicTrans2	ai4bharat/indictrans2-en-indic-1B	Native (Indic-trained)	Lowest
NLLB-200	facebook/nllb-200-distilled-600M	Strong (multilingual)	Low
Helsinki-dra	Helsinki-NLP/opus-mt-en-dra	None (romanised only)	Misleading*
mBART-50	facebook/mbart-large-50-many-to-many-mmt	ta_IN tag	Moderate
MADLAD-400	google/madlad400-3b-mt	Strong (400-lang)	Low

* Helsinki-dra's low expansion is an artifact of mapping all Tamil Unicode to <unk>, not genuine compression.

Part A – Batch Translation & Evaluation

IndicTrans2 was evaluated on 500 English-Tamil sentence pairs drawn from the Helsinki-NLP/opus-100 en-ta test split. The model's IndicProcessor pipeline is load-bearing infrastructure: it normalises scripts, injects language tags (eng_Latn → tam_TamI), and cleans punctuation before any tokens are generated. Without it, even correctly generated token sequences render as garbled output.

Key Findings

1. BLEU distribution is right-skewed: Short, structurally simple sentences score well. Complex subordinate clauses and idiomatic expressions introduce n-gram divergence that BLEU penalises heavily, even when the translation is semantically correct.

2. BLEU understates real quality: Tamil allows significant word-order flexibility and multiple valid surface forms for the same meaning. A model that produces a correct translation in a different valid phrasing from the reference gets penalised. BLEU is a surface metric, not a meaning metric.

3. IndicProcessor is non-negotiable: The postprocessing step removes BPE boundary markers, normalises Tamil punctuation, and cleans whitespace artifacts. The model and its processor are a system; separating them breaks both.

Additional Metrics

Metric	Corpus Score	Avg Sentence	Max	Min	Note
chrF	40.43	40.47	100.0	0.0	Higher is better
chrF++	35.34	36.64	100.0	0.0	Higher is better
TER	88.53	N/A	N/A	N/A	Lower is better
ROUGE-L	5.69	5.69 (avg)	100.0	0.0	Higher is better; sentence avg

4. chrF (40.43) substantially outperforms BLEU (7.28): This is expected for Tamil. BLEU penalises any surface-form mismatch including morphological variants that are semantically equivalent. chrF's character-level overlap is far more forgiving of valid Tamil morphological alternatives, making it the more reliable metric for this language pair.

5. TER (88.53) is high due to Tamil's SOV word order: TER counts edits needed to convert a hypothesis into the reference. Tamil is SOV while English is SVO, so even a correct translation requires many reorderings relative to the reference, inflating TER. This is expected and does not indicate poor performance.

6. ROUGE-L (5.69) reflects the same word-level limitations as BLEU: ROUGE-L measures longest common subsequence at word level. For Tamil, a single agglutinated word can correspond to an entire English phrase, so word-level LCS rarely aligns even when the translation is semantically correct. The rouge-score library computes only sentence-level scores; the corpus value reported is the sentence average, consistent with the original ROUGE paper.

Part B – Token-Based Comparative Analysis

All five models translated 15 English sentences across three complexity tiers. Six metrics were computed per sentence per model: source/target token counts, expansion ratio, average word length, subword fragmentation, and unknown token rate. The results form a clear hierarchy that mirrors the tokenizer vocabulary quality seen in Part C.

Architectural Fixes

- Helsinki-dra requires >>tam<< prefix per sentence or it defaults to a non-Tamil Dravidian output.
- IndicTrans2 tokenizer raises AssertionError on plain text; fixed with word-count fallback for token metrics.
- mBART-50 requires explicit `src_lang = "en_XX"` and `forced_bos_token_id = lang_code_to_id["ta_IN"]`.

Key Findings

1. IndicTrans2 achieves the lowest expansion ratio: Its Indic-trained SentencePiece vocabulary encodes Tamil agglutinative words in 3-4 tokens where general multilingual models need 8-12. This is not luck; it is a vocabulary design decision.

2. Helsinki-dra's unk_rate approaches 1.0: Trained on romanised Dravidian text, it has zero Tamil Unicode coverage. Its deceptively low expansion ratio is an artifact of token collapse; every Tamil word becomes <unk>.

3. Complexity drives fragmentation variance: Simple sentences (Set 1) show tight, consistent ratios across models. Complex sentences (Set 3) show fragmentation spikes, especially for models with weak Tamil vocabulary coverage.

Part C – Indic Token Behavior Analysis

No models. No translations. Just tokenizers and raw Tamil Unicode. 20 Tamil sentences spanning education, science, healthcare, and culture were fed directly to each tokenizer. The question: can the tokenizer even see Tamil? The answer varies dramatically.

Tokenization of மருத்துவமனையில் ("in the hospital", 16 chars)

Tokenizer	Tokens for: மருத்துவமனையில்	Count
IndicTrans2	—மருத்து வமனை யில் ஂ	~4
NLLB-200	—மருத் துவ மனை யி ல் ஂ	~6
mBART-50	—மருத்து வமனை யில்	~3
MADLAD-400	—மருத் துவ மனை யில்	~4
Helsinki-dra	— <unk>	2 (failure)

Key Findings

- Vocabulary is destiny:** Models trained with Tamil Unicode in their SentencePiece vocab achieve meaningful morpheme-level segmentation. Models without it produce <unk>. No model size or architecture compensates for a vocabulary that never saw Tamil script.
- Unicode combining marks are the hidden trap:** Tamil vowel signs and virama (U+0BCD) visually attach to consonants but can be split by naive tokenizers into isolated 1-character tokens, destroying syllable structure and wasting attention heads.
- Token count has a quadratic cost:** Every additional token increases Transformer attention computation quadratically. IndicTrans2's compact Tamil encoding is not just elegant; it is economically significant at inference scale.

Conclusion: Across all three parts, a single thesis holds: tokenizer vocabulary design is the decisive factor in Tamil NLP quality. IndicTrans2's Indic-first SentencePiece model gives it a structural advantage no general multilingual model can overcome at the vocabulary layer. NLLB-200 and MADLAD-400 are strong alternatives; Helsinki-dra is a cautionary tale. The low BLEU (7.28), ROUGE-L (5.69), and high TER (88.53) in Part A should not be read as evidence of poor translation; all three are ill-suited to Tamil's agglutinative, SOV structure. chrF (40.43) is the most appropriate automatic metric for this language pair. The token analyses in Parts B and C are not academic exercises; they directly explain the quality rankings in Part A. In Tamil NLP, the vocabulary is the model.